

Getting Started

The goal of this assignment is to develop a method to predict the electricity load demand of 3 individual users. For each user, we are given following 2 datasets which cover 1 calendar year:

- Energy usage history (in kW) with 30-minute or 1-minute interval
- Weather history with 1-hour interval

Our objective is to create models that predict the future electricity consumption of these customers and measure the accuracy of your predictions by the Mean Absolute Error.

First and foremost, lets import the libraries and & datasets.

```
In [1]: import pandas as pd
import numpy as np
import datetime
from datetime import timedelta
from IPython.core.display import display
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error
from statsmodels.tsa.arima_model import ARIMA
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.ensemble import RandomForestRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import ShuffleSplit
from sklearn import preprocessing

# Loading the data from csv files

energyB = pd.read_csv('./data/HomeB-meter1_2014.csv')
energyC = pd.read_csv('./data/HomeC-meter1_2016.csv')
energyF = pd.read_csv('./data/HomeF-meter3_2016.csv')
weatherB = pd.read_csv('./data/homeB2014.csv')
weatherC = pd.read_csv('./data/homeC2016.csv')
weatherF = pd.read_csv('./data/homeF2016.csv')
```

Data Exploration

```
In [2]: display(weatherB.head())
display(energyB)
```

	temperature	icon	humidity	visibility	summary	apparentTemperature	pressure	windSpeed
0	34.98	partly-cloudy-night	0.64	10.00	Partly Cloudy	28.62	1017.69	7.75
1	16.49	clear-night	0.62	10.00	Clear	16.49	1022.76	2.71
2	14.63	clear-night	0.68	10.00	Clear	6.87	1022.32	4.84
3	13.31	clear-night	0.71	10.00	Clear	6.49	1021.64	4.00
4	13.57	clear-night	0.71	9.93	Clear	7.29	1020.73	3.67

	Date & Time	use [kW]	gen [kW]	Grid [kW]	AC [kW]	Furnace [kW]	Cellar Lights [kW]	Washer [kW]	First Floor lights [kW]	Utility Ba: Ba
0	2014-01-01 00:00:00	0.304439	0.0	0.304439	0.000058	0.009531	0.005336	0.000126	0.011175	0.
1	2014-01-01 00:30:00	0.656771	0.0	0.656771	0.001534	0.364338	0.005522	0.000043	0.003514	0.
2	2014-01-01 01:00:00	0.612895	0.0	0.612895	0.001847	0.417989	0.005504	0.000044	0.003528	0.
3	2014-01-01 01:30:00	0.683979	0.0	0.683979	0.001744	0.410653	0.005556	0.000059	0.003499	0.
4	2014-01-01 02:00:00	0.197809	0.0	0.197809	0.000030	0.017152	0.005302	0.000119	0.003694	0.
...
17515	2014-12-31 21:30:00	1.560890	0.0	1.560890	0.003226	0.392996	0.006342	0.000872	0.030453	0.
17516	2014-12-31 22:00:00	0.958447	0.0	0.958447	0.000827	0.027369	0.006326	0.000811	0.030391	0.
17517	2014-12-31 22:30:00	0.834462	0.0	0.834462	0.001438	0.170561	0.020708	0.000636	0.012631	0.

	Date & Time	use [kW]	gen [kW]	Grid [kW]	AC [kW]	Furnace [kW]	Cellar Lights [kW]	Washer [kW]	First Floor lights [kW]	Uti Ba: Ba
17518	2014-12-31 23:00:00	0.543863	0.0	0.543863	0.001164	0.153533	0.008423	0.000553	0.003832	0.
17519	2014-12-31 23:30:00	0.414441	0.0	0.414441	0.000276	0.009223	0.006619	0.000526	0.003818	0.

17520 rows × 18 columns

In [3]: energyB.describe()

Out[3]:

	use [kW]	gen [kW]	Grid [kW]	AC [kW]	Furnace [kW]	Cellar Lights [kW]	Washer [kW]
count	17520.000000	17520.0	17520.000000	17520.000000	17520.000000	17520.000000	17520.00000
mean	0.662905	0.0	0.662905	0.088999	0.085888	0.011036	0.00306
std	0.678399	0.0	0.678399	0.438887	0.129054	0.013123	0.02044
min	0.011083	0.0	0.011083	0.000000	0.000117	0.000083	0.00000
25%	0.314125	0.0	0.314125	0.000030	0.009340	0.005414	0.00009
50%	0.468725	0.0	0.468725	0.000069	0.009704	0.005881	0.00021
75%	0.700617	0.0	0.700617	0.000707	0.143531	0.007042	0.00033
max	6.833205	0.0	6.833205	3.687768	0.437212	0.146692	0.81916

In [4]: weatherB.describe()

Out[4]:

	temperature	humidity	visibility	apparentTemperature	pressure	windSpeed	cloudCover
count	8760.000000	8760.000000	8760.000000	8760.000000	8760.000000	8760.000000	729
mean	48.062076	0.682888	9.025791	45.289160	1016.450749	6.534568	
std	19.694743	0.188763	1.859263	22.860668	7.903670	3.884500	
min	-10.070000	0.140000	0.320000	-18.280000	979.980000	0.030000	
25%	33.165000	0.530000	9.040000	27.967500	1011.530000	3.630000	
50%	49.220000	0.710000	9.970000	47.360000	1016.430000	5.850000	
75%	63.832500	0.860000	10.000000	63.832500	1021.310000	8.692500	
max	89.460000	0.960000	10.000000	97.520000	1042.400000	24.750000	

We picked up House B data to analyse it further. The weather data set has different columns

specifying the parameters and their corresponding values.

We are particularly interested in the count.

For house B we have **8760 data points** on weather which correspond to hourly data for a year. (**365 days * 24**). While, for energy usage , we have 17520 data points which correspond to bihourly data points for the year. (365 days * 24 hours * 2 (half hour)).

We need to normalise the data to bring it in the same order and we shall address these in the next section.

```
In [5]: energyC.tail()
```

Out[5]:

	Date & Time	use [kW]	gen [kW]	House overall [kW]	Dishwasher [kW]	Furnace 1 [kW]	Furnace 2 [kW]	Home office [kW]	Fridge [kW]
503905	2016-12-15 22:25:00	1.601233	0.003183	1.601233	0.000050	0.085267	0.642417	0.041783	0.005267
503906	2016-12-15 22:26:00	1.599333	0.003233	1.599333	0.000050	0.104017	0.625033	0.041750	0.005233
503907	2016-12-15 22:27:00	1.924267	0.003217	1.924267	0.000033	0.422383	0.637733	0.042033	0.004983
503908	2016-12-15 22:28:00	1.978200	0.003217	1.978200	0.000050	0.495667	0.620367	0.042100	0.005333
503909	2016-12-15 22:29:00	1.990950	0.003233	1.990950	0.000050	0.494700	0.634133	0.042100	0.004917

```
In [6]: energyF.tail()
```

Out[6]:

	Date & Time	Usage [kW]	Generation [kW]	Net_Meter [kW]	Volt [kW]	Garage_E [kW]	Garage_W [kW]	Phase_A [kW]	Phase_B [kW]
503920	2016-12-15 22:40:00	0.643783	0.009100	0.652883	0.002200	0.000000	0.000350	0.490050	0.153750
503921	2016-12-15 22:41:00	1.135383	0.009117	1.144500	0.002200	0.000000	0.000350	0.492050	0.643750
503922	2016-12-15 22:42:00	1.395117	0.008150	1.403267	0.002200	0.000017	0.000350	0.427983	0.967250
503923	2016-12-15 22:43:00	0.624050	0.007933	0.631983	0.002200	0.000000	0.000367	0.474950	0.143750
503924	2016-12-15 22:44:00	0.597250	0.006600	0.603850	0.002167	0.000000	0.000350	0.454467	0.143750

For House B we have energy usage data till 31st December, where as for House C & F, we have data points only till 15th Dec. Hence, as per assignment specification, as we are using data till 30th November as training set for House B and till 15th Nov for House C and F. This gives us effectively roughly 30days datapoints for testing dataset for the predictions.

Data Preprocessing

Before data can be used as input for machine learning algorithms, it often must be cleaned, formatted, and restructured — this is typically known as **preprocessing**. There are some qualities about certain features that must be adjusted before we go ahead. This preprocessing can help tremendously with the outcome and predictive power of nearly all learning algorithms.

For the purposes of this project, the following preprocessing steps have been made to the dataset:

1. The column 'Date & Time' needs to be converted to datetime data type. This would help in operations over the data.

```
In [7]: # Convert to Date Time Format
energyB['Date & Time'] = pd.to_datetime(energyB['Date & Time'])
energyC['Date & Time'] = pd.to_datetime(energyC['Date & Time'])
energyF['Date & Time'] = pd.to_datetime(energyF['Date & Time'])
```

2. We need to change the column name of energyF data frame from 'Usage [kW]' to 'use [kW]' for consistency across different datasets. This would help in accessing the datasets in uniform manner for referring the columns.

```
In [8]: # Changing column name
energyF = energyF.rename(columns={'Usage [kW]': 'use [kW]'})
```

3. Cardinality of Datasets

As observed above, we need to bring the features and target datasets to same order so as to map them. Lets check the shape of the data.

```
In [9]: print("Shape of Data for House B (Energy Usage):", energyB.shape)
print("Shape of Data for House B (Weather):", weatherB.shape)
print("Shape of Data for House C (Energy Usage):", energyC.shape)
print("Shape of Data for House C (Weather):", weatherC.shape)
print("Shape of Data for House F (Energy Usage):", energyF.shape)
print("Shape of Data for House F (Weather):", weatherF.shape)
```

```
Shape of Data for House B (Energy Usage): (17520, 18)
Shape of Data for House B (Weather): (8760, 14)
Shape of Data for House C (Energy Usage): (503910, 19)
Shape of Data for House C (Weather): (8760, 14)
Shape of Data for House F (Energy Usage): (503925, 10)
Shape of Data for House F (Weather): (8760, 14)
```

Hence we see that the weather data available is hourly data,(i.e. number of rows = 365 * 24) but the

energy data is bihourly for House C and Minute Level for House C & F.

Preparing Hourly and Daily timeline datasets

The count of the above datasets show. We need to convert the number of rows in Weather data corresponds to number of hours in a year ($365 \times 24 = 8760$).

Hourly

To prepare Hously data on energy usage, we need to merge the bihourly (House B)or minute level (House C & F) datapoints.

Daily

To convert to daily data, we round the dates to their corresponding days and then group by on date.

For **Energy usage data**, we take the **sum** of each cell values, since energy usage for hour would be sum of usages of every minute in that hour.

For **weather data** values, we take the **median on group by instead of mean, because it isn't influenced by extremely large values or outliers** . With these basis we group the data and prepare the features and target datasets.

```
In [10]: # Merge energy bihourly data into hourly data
energyB['Date & Time'] = energyB["Date & Time"].dt.floor('H')
energyC['Date & Time'] = energyC["Date & Time"].dt.floor('H')
energyF['Date & Time'] = energyF["Date & Time"].dt.floor('H')
energy_hourly_B = energyB.groupby('Date & Time').sum().reset_index()
energy_hourly_C = energyC.groupby('Date & Time').sum().reset_index()
energy_hourly_F = energyF.groupby('Date & Time').sum().reset_index()
```

```
In [11]: energyB['Date & Time'] = energyB["Date & Time"].dt.floor('D')
energyC['Date & Time'] = energyC["Date & Time"].dt.floor('D')
energyF['Date & Time'] = energyF["Date & Time"].dt.floor('D')
energy_daily_B = energyB.groupby('Date & Time').sum().reset_index()
energy_daily_C = energyC.groupby('Date & Time').sum().reset_index()
energy_daily_F = energyF.groupby('Date & Time').sum().reset_index()
```

```

In [12]: def getDateTIme(inputYear, df):
    currenttime = datetime.datetime(inputYear,1,1,0,0,0)
    datearr = []
    for index, row in df.iterrows():
        datearr.append(currenttime)
        #add delta of one hour to each iteration
        currenttime = currenttime + timedelta(hours=1)
    return datearr

weatherB['Date & Time'] = getDateTIme(2014, weatherB)
weatherC['Date & Time'] = getDateTIme(2016, weatherC)
weatherF['Date & Time'] = getDateTIme(2016, weatherF)
weather_hourly_B = weatherB.copy()
weather_hourly_C = weatherC.copy()
weather_hourly_F = weatherF.copy()

weatherB['Date & Time'] = weatherB["Date & Time"].dt.floor('D')
weatherC['Date & Time'] = weatherC["Date & Time"].dt.floor('D')
weatherF['Date & Time'] = weatherF["Date & Time"].dt.floor('D')
weather_daily_B = weatherB.groupby('Date & Time').median().reset_index()
weather_daily_C = weatherC.groupby('Date & Time').median().reset_index()
weather_daily_F = weatherF.groupby('Date & Time').median().reset_index()

In [13]: weatherData=[weather_hourly_B, weather_hourly_C, weather_hourly_F, weather_
energyData = [energy_hourly_B, energy_hourly_C, energy_hourly_F, energy_dai
tag = ['House B (Hourly)', 'House C (Hourly)', 'House F (Hourly)', 'House B (D

models=[]
maeBHourly=[]
maeCHourly=[]
maeFHourly=[]
maeBDaily=[]
maeCDaily=[]
maeFDaily=[]
aggregate=[maeBHourly,maeCHourly,maeFHourly,maeBDaily,maeCDaily,maeFDaily]


```

Data Visualisation


Let's visualise the different aggregates of the energy data and see what insights we can get from it.

Hourly Plots


```
In [14]: # Plot graphs for both the diagrams
energy_hourly_B.plot(kind='line', x='Date & Time', y='use [kW]', figsize=(3
plt.show())
```



```
In [15]: energy_hourly_C.plot(kind='line', x='Date & Time', y='use [kW]', figsize=(3
plt.show())
```




```
In [16]: energy_hourly_F.plot(kind='line', x='Date & Time', y='use [kW]', figsize=(3
plt.show())
```




Daily Plots


```
In [17]: energy_daily_B.plot(kind='line', x='Date & Time', y='use [kW]', figsize=(30
plt.show())
```




```
In [18]: energy_daily_C.plot(kind='line', x='Date & Time', y='use [kW]', figsize=(30, 10),  
plt.show())
```



```
In [19]: energy_daily_F.plot(kind='line', x='Date & Time', y='use [kW]', figsize=(30, 10),  
plt.show())
```



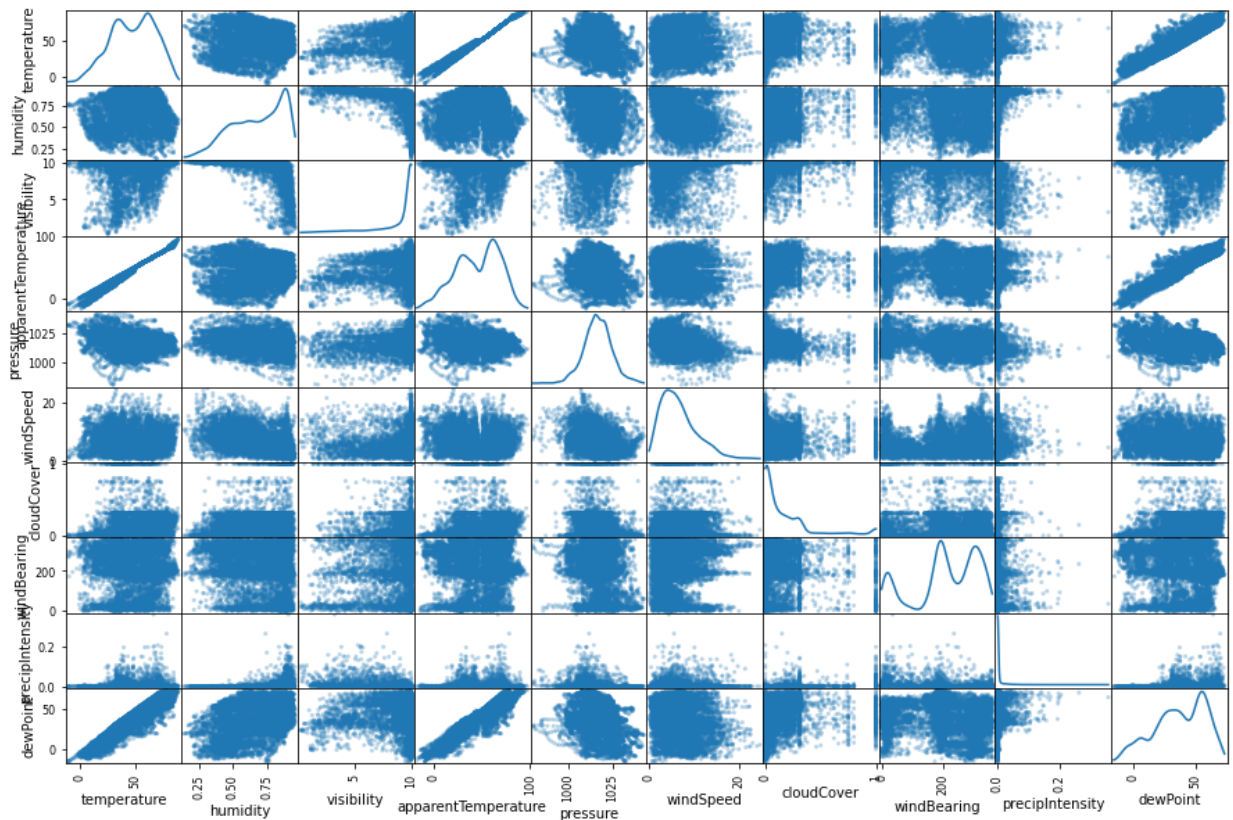
From the above visualisations we can clearly see, as for House B & C , we observe significant spike during the late summer months. This might be due to excessive usage of cooling appliances. House F on the other hand has high energy usage (High avg mean) all the year round. Probably House F would be a data of a bigger house or public place or any kind of hotels.

Visualizing Feature Distributions

To get a better understanding of the dataset, we can construct a scatter matrix of each numerical features present in the weather dataset. If we believe that any feature is correlated with another features, it might be relevant for identifying the energy demand and the scatter matrix might show a correlation between that feature and another feature in the dataset.


```
In [20]: from pandas.plotting import scatter_matrix

data = weather_hourly_B[['temperature', 'humidity', 'visibility', 'apparentTem
scatter_matrix(data, alpha = 0.3, figsize = (15,10), diagonal = 'kde');
```



From the above diagram, from the data points in the scatter plot we can clearly see that **Temperature, Apparent Temperature and Dew point** are the most tightly clustered along an imaginary line, hence they have some positive correlation with each other. To further accentuate our observation, let's compute pairwise correlation of columns and plot it on heat map.

```
In [21]: feature_corr = data.corr()
fig, ax = plt.subplots(figsize=(20,12))
sns.heatmap(feature_corr, annot=True, linewidths=.5)
plt.show()
```



The high correlation values of **Temperature**, **Apparent Temperature** and **Dew point** shows that they are closely amongst themselves. Hence as part of training & prediction, we can use amongst these to train and test models on the dataset. Since apparent temperature has a close synergy with Temperature feature, we shall use **temperature and dewPoint** to train our models.

```
In [22]: linkedFeatures = ['temperature', 'dewPoint']
```

Now lets try and create models to predict the future electricity consumption of these houses on hourly and daily data points.

Training & Prediction

Now that our data-preprocessing is done, we will use the datasets to train models, predict the energy usage and calculate the accuracy in terms of Mean Absolute Error.

We are using techniques like **Naive Method**, **Linear Regression**, **Decision tree Regression**, **Random Forest** and **ARIMA Model** to produce forecasts. For parameter tuning, I am using **Grid Search** Method to tune the hyperparameters.

As per the requirement specification, I am splitting the data set into two parts based on the date. I am using Jan to Nov data as the training set and Dec datapoints to test and calculate the MAE.

Naive Method Prediction

Naive method is an estimating technique in which the last period's actual values are used to predict the future forecast, without adjusting them or attempting to establish causal factors. It is here used here as a baseline to compare forecasts generated by the better (sophisticated) techniques.

As one can see below, we are preparing test and train dataset based on split on date. For House B, we use training data from 1st Jan to 2014-12-01 00:00:00. For House C & F, We use training data from 1st Jan to "2016-11-15 00:00:00", since this would gave 30days data to test in both the cases. Since we want to compare amongst similar data, we are storing the split points and using the same splits to test other algorithms so that we can compare them with Naive method.

```
In [23]: # NAIVE Method Prediction
def naiveMethodPredictor(dataset, split):
    train_Y = dataset[:split]
    test_Y = dataset[split:]
    count = len(test_Y)
    predict_Y = pd.DataFrame(np.asarray(train_Y)[len(train_Y) - 1][0], index=test_Y.index)
    mae = mean_absolute_error(test_Y, predict_Y)
    return mae
```

```
In [24]: models.append("Naive Method")

splitHB = len(energy_hourly_B.loc[energy_hourly_B['Date & Time'] <= '2014-1
maeVal = naiveMethodPredictor(energy_hourly_B[['use [kW]']], splitHB)
maeBHourly.append(maeVal)
print("The Mean Absolute Error of Naive method for Hourly points on House B

splitHC = len(energy_hourly_C.loc[energy_hourly_C['Date & Time'] <= '2016-1
maeVal = naiveMethodPredictor(energy_hourly_C[['use [kW]']], splitHC)
maeCHourly.append(maeVal)
print("The Mean Absolute Error of Naive method for Hourly points on House C

splitHF = len(energy_hourly_F.loc[energy_hourly_F['Date & Time'] <= '2016-1
maeVal = naiveMethodPredictor(energy_hourly_F[['use [kW]']], splitHF)
maeFHourly.append(maeVal)
print("The Mean Absolute Error of Naive method for Hourly points on House F

splitDB = len(energy_daily_B.loc[energy_daily_B['Date & Time'] <= '2014-12-
maeVal = naiveMethodPredictor(energy_daily_B[['use [kW]']], splitDB)
maeBDaily.append(maeVal)
print("The Mean Absolute Error of Naive method for Daily points on House B

splitDC = len(energy_daily_C.loc[energy_daily_C['Date & Time'] <= '2016-11-
maeVal = naiveMethodPredictor(energy_daily_C[['use [kW]']], splitDC)
maeCDaily.append(maeVal)
print("The Mean Absolute Error of Naive method for Daily points on House C

splitDF = len(energy_daily_F.loc[energy_daily_F['Date & Time'] <= '2016-11-
maeVal = naiveMethodPredictor(energy_daily_F[['use [kW]']], splitDF)
maeFDaily.append(maeVal)
print("The Mean Absolute Error of Naive method for Daily points on House F
```

```
The Mean Absolute Error of Naive method for Hourly points on House B is :
0.4795294990740242
The Mean Absolute Error of Naive method for Hourly points on House C is :
25.24953158025876
The Mean Absolute Error of Naive method for Hourly points on House F is :
47.94122041828032
The Mean Absolute Error of Naive method for Daily points on House B is :
6.231544518466666
The Mean Absolute Error of Naive method for Daily points on House C is :
252.46967000100022
The Mean Absolute Error of Naive method for Daily points on House F is :
440.9711255574666
```

```
In [25]: splitPoints=[]
splitPoints.append(splitHB)
splitPoints.append(splitHC)
splitPoints.append(splitHF)
splitPoints.append(splitDB)
splitPoints.append(splitDC)
splitPoints.append(splitDF)
```

In an attempt to compare apples to apples, we are memoizing the split points and use the same to test and train different models.

```
In [26]: from matplotlib.pyplot import figure
figure(num=None, figsize=(10, 30), dpi=200, facecolor='w', edgecolor='k')

def plotvalues(x, y1, y2, house, model1, model2, interval):
    plt.title('Time Series plot at: ' + house + " " + interval)
    plt.xlabel('Date & Time')
    plt.ylabel('Use [kWh]')
    plt.plot(x, y1, 'b', label=model1)
    plt.plot(x, y2, 'r', label=model2)
    plt.legend(loc='upper left')
    plt.show()
```

<Figure size 2000x6000 with 0 Axes>

Linear Regression

Linear regression attempts to model the relationship between two variables by fitting a linear equation to observed data. One variable is considered to be an explanatory variable, and the other is considered to be a dependent variable. We are using the the features which we obtained from feature distribution. I am using the split points I extracted above to prepare training and testing datasets.

For training features, I am using the temperature and dewpoint for training and predict energy usages based on it. Then I have plots of graphs for the actual energy usage values vs the predicted values.

Normalising the features

Since we will be using the features like Temperature and Dew Point, the values of these features are in different magnitude. For example for House B weather data, as seen above, (temperature has max value of 89 and min value of -10) and dew point varies from (-15 to 72). Hence we are normalising the features dataset to bring them into same order.

In [27]: *# Train Linear Regression model and predict*

```
def LinearRegressionModel(features, target, split, house, interval):
    df = features[linkedFeatures].reset_index(drop=True)
    min_max_scaler = preprocessing.MinMaxScaler()
    scaled = min_max_scaler.fit_transform(df.values)
    scaled = pd.DataFrame(scaled, columns=df.columns, index=df.index)
    xtrain = scaled[:split].reset_index(drop=True)

    ytrain = target[:split][['use [kW]']].reset_index(drop=True)
    ytest = target[split:][['use [kW]']].reset_index(drop=True)
    xtest = scaled[split : (split + len(ytest))].reset_index(drop=True)
    linear_regressor = LinearRegression()
    model = linear_regressor.fit(xtrain, ytrain)
    ypredict = linear_regressor.predict(xtest)
    mae = mean_absolute_error(ytest, ypredict)
    plotvalues(features[split : (split + len(ytest))][["Date & Time"]], yte
    return mae
```

```

In [28]: models.append("Linear Regression")

maeVal = LinearRegressionModel(weather_hourly_B, energy_hourly_B, splitHB,
maeBHourly.append(maeVal)
print("The Mean Absolute Error for Linear Regression for hourly points on H

maeVal = LinearRegressionModel(weather_hourly_C, energy_hourly_C, splitHC,
maeCHourly.append(maeVal)
print("The Mean Absolute Error for Linear Regression for hourly points on H


maeVal = LinearRegressionModel(weather_hourly_F, energy_hourly_F, splitHF,
maeFHourly.append(maeVal)
print("The Mean Absolute Error for Linear Regression for hourly points on H

maeVal = LinearRegressionModel(weather_daily_B, energy_daily_B, splitDB, "H
maeBDaily.append(maeVal)
print("The Mean Absolute Error for Linear Regression for daily points on Ho


maeVal = LinearRegressionModel(weather_daily_C, energy_daily_C, splitDC, "H
maeCDaily.append(maeVal)
print("The Mean Absolute Error for Linear Regression for daily points on Ho

maeVal = LinearRegressionModel(weather_daily_F, energy_daily_F, splitDF, "H
maeFDaily.append(maeVal)
print("The Mean Absolute Error for Linear Regression for daily points on Ho


```




The Mean Absolute Error for Linear Regression for hourly points on House B is 0.4637947126337114



The Mean Absolute Error for Linear Regression for hourly points on House C is 21.191972179063775




The Mean Absolute Error for Linear Regression for hourly points on House F is 45.50591348229302




The Mean Absolute Error for Linear Regression for daily points on House B

```
is 6.157594756866734
```



The Mean Absolute Error for Linear Regression for daily points on House C is 266.6723313494483



The Mean Absolute Error for Linear Regression for daily points on House F is 431.7812314287284

Decision Tree Regressor

A decision tree is a supervised machine learning model used to predict a target by learning decision rules from features. As the name suggests, we can think of this model as breaking down

our data by making a decision based on asking a series of questions. Initially we split the data based on the split points to prepare training and testing datasets and then we train models to get predicted values. We run it on default parameters in the first attempt. Following that, we try to improve the MAE by hyperparameter tuning using Grid Search as shown below.

```
In [29]: def DecisionTree(features, target, split):  
    xtrain = features[:split][linkedFeatures].reset_index(drop=True)  
    ytrain = target[:split][['use [kW]']].reset_index(drop=True)  
    ytest = target[split:][['use [kW]']].reset_index(drop=True)  
    xtest = features[split : (split + len(ytest))][linkedFeatures].reset_in  
    regressor = DecisionTreeRegressor(random_state=0)  
    model = regressor.fit(xtrain, ytrain.values.ravel())  
    mae = mean_absolute_error(ytest, model.predict(xtest))  
    return mae
```

Hourly & Daily MAE for Houses B, C, & F

```
In [30]: models.append("Decision Tree (Without Grid Search)")

maeVal = DecisionTree(weather_hourly_B, energy_hourly_B, splitHB)
maeBHourly.append(maeVal)
print("The Mean Absolute Error for DecisionTree Model for hourly points on House B is", maeVal)

maeVal = DecisionTree(weather_hourly_C, energy_hourly_C, splitHC)
maeCHourly.append(maeVal)
print("The Mean Absolute Error for DecisionTree Model for hourly points on House C is", maeVal)

maeVal = DecisionTree(weather_hourly_F, energy_hourly_F, splitHF)
maeFHourly.append(maeVal)
print("The Mean Absolute Error for DecisionTree Model for hourly points on House F is", maeVal)

maeVal = DecisionTree(weather_daily_B, energy_daily_B, splitDB)
maeBDaily.append(maeVal)
print("The Mean Absolute Error for DecisionTree Model for daily points on House B is", maeVal)

maeVal = DecisionTree(weather_daily_C, energy_daily_C, splitDC)
maeCDaily.append(maeVal)
print("The Mean Absolute Error for DecisionTree Model for daily points on House C is", maeVal)

maeVal = DecisionTree(weather_daily_F, energy_daily_F, splitDF)
maeFDaily.append(maeVal)
print("The Mean Absolute Error for DecisionTree Model for daily points on House F is", maeVal)
```

```
The Mean Absolute Error for DecisionTree Model for hourly points on House B is 0.6952422633539705
The Mean Absolute Error for DecisionTree Model for hourly points on House C is 31.51969418227089
The Mean Absolute Error for DecisionTree Model for hourly points on House F is 63.34120420022507
The Mean Absolute Error for DecisionTree Model for daily points on House B is 9.442799666933333
The Mean Absolute Error for DecisionTree Model for daily points on House C is 333.6515294442667
The Mean Absolute Error for DecisionTree Model for daily points on House F is 430.68034778346725
```

Grid Search | Parameter Tuning & Optimising Model Performance

Grid search is an algorithm with which we tune hyperparameters (example `max_depth` in our case) of our model. The input to grid search is possible values of hyperparameters and possible tuning metric.

The algo tries all possible combinations of hyperparameters in a grid and evaluates the performance of each combo with some cross validation set. The output is the hyperparameter combo which produces the best result.

The performance metric is required to assess the which is the best performing hyperparameters to be used in learning algorithm. The grid signifies the exhaustive nature of approach to try out all possible combinations of the hyperparameters.

We are using **R2 Score**, to quantify your model's performance. The coefficient of determination for a model is a useful statistic in regression analysis, as it often describes how "good" that model is at making predictions.

```
In [31]: from sklearn.metrics import make_scorer
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import r2_score

def performance_metric(y_true, y_predict):
    """ Calculates and returns the performance score between
        true and predicted values based on the metric chosen. """
    score = r2_score(y_true, y_predict)
    return score

def fit_model(X, y):
    """ Performs grid search over the 'max_depth' parameter for a
        decision tree regressor trained on the input data [X, y]. """
    regressor = DecisionTreeRegressor()
    params = dict(max_depth=[1,2,3,4,5,6,7,8,9,10])
    scoring_fnc = make_scorer(performance_metric)
    grid = GridSearchCV(regressor, params, scoring=scoring_fnc)
    grid = grid.fit(X, y)
    return grid.best_estimator_

def DecisionTreeModel(features, target, split, house):
    xtrain = features[:split][linkedFeatures].reset_index(drop=True)
    ytrain = target[:split][['use [kW]']].reset_index(drop=True)
    ytest = target[split:][['use [kW]']].reset_index(drop=True)
    xtest = features[split : (split + len(ytest))][linkedFeatures].reset_in
    model = fit_model(xtrain, ytrain)
    print("Parameter 'max_depth' is {} for the optimal model.".format(model
    ypredict = model.predict(xtest)
    plotvalues(features[split : (split + len(ytest))][["Date & Time"]], yte
    mae = mean_absolute_error(ytest, ypredict)
    return mae
```

Hourly and Daily datapoints

```
In [32]: models.append("Decision Tree (With Grid Search)")

maeVal = DecisionTreeModel(weather_hourly_B, energy_hourly_B, splitHB, "House B")
maeBHourly.append(maeVal)
print("The Mean Absolute Error for DecisionTree Model for hourly points on House B is: ", maeVal)

maeVal = DecisionTreeModel(weather_hourly_C, energy_hourly_C, splitHC, "House C")
maeCHourly.append(maeVal)
print("The Mean Absolute Error for DecisionTree Model for hourly points on House C is: ", maeVal)


maeVal = DecisionTreeModel(weather_hourly_F, energy_hourly_F, splitHF, "House F")
maeFHourly.append(maeVal)
print("The Mean Absolute Error for DecisionTree Model for hourly points on House F is: ", maeVal)

maeVal = DecisionTreeModel(weather_daily_B, energy_daily_B, splitDB, "House B")
maeBDaily.append(maeVal)
print("The Mean Absolute Error for DecisionTree Model for daily points on House B is: ", maeVal)


maeVal = DecisionTreeModel(weather_daily_C, energy_daily_C, splitDC, "House C")
maeCDaily.append(maeVal)
print("The Mean Absolute Error for DecisionTree Model for daily points on House C is: ", maeVal)

maeVal = DecisionTreeModel(weather_daily_F, energy_daily_F, splitDF, "House F")
maeFDaily.append(maeVal)
print("The Mean Absolute Error for DecisionTree Model for daily points on House F is: ", maeVal)
```


Parameter 'max_depth' is 2 for the optimal model.




The Mean Absolute Error for DecisionTree Model for hourly points on House B is 0.48084775861604884
Parameter 'max_depth' is 4 for the optimal model.



The Mean Absolute Error for DecisionTree Model for hourly points on House C is 21.50087723226878
Parameter 'max_depth' is 2 for the optimal model.




The Mean Absolute Error for DecisionTree Model for hourly points on House F is 50.25079070120949
Parameter 'max_depth' is 1 for the optimal model.




The Mean Absolute Error for DecisionTree Model for daily points on House B is 6.228120276068112

Parameter 'max_depth' is 4 for the optimal model.



The Mean Absolute Error for DecisionTree Model for daily points on House C is 201.46049513978446

Parameter 'max_depth' is 1 for the optimal model.



The Mean Absolute Error for DecisionTree Model for daily points on House F is 424.30888697181285

```
In [33]: decisionTree = pd.DataFrame({
    'Models': models,
    'House B (Hourly)': maeBHourly,
    'House C (Hourly)': maeCHourly,
    'House F (Hourly)': maeFHourly,
    'House B (Daily)': maeBDaily,
    'House C (Daily)': maeCDaily,
    'House F (Daily)': maeFDaily,
})
print("Differences between MAEs:")
display(decisionTree.loc[2:3])
```

Differences between MAEs:

	Models	House B (Hourly)	House C (Hourly)	House F (Hourly)	House B (Daily)	House C (Daily)	House F (Daily)
2	Decision Tree (Without Grid Search)	0.695242	31.519694	63.341204	9.44280	333.651529	430.680348
3	Decision Tree (With Grid Search)	0.480848	21.500877	50.250791	6.22812	201.460495	424.308887

Hence, its clear quantitatively, by how much **Grid Search parameter Tuning** improves the MAE of the predictions made by different datasets.

Random Forest

Random Forest Regression is a supervised learning algorithm that uses ensemble learning method for regression. As we saw above the benefits of parameter tuning, we take multiple max depth values as input for tuning parameters. Then using the best estimator to calculate the MAE.

```
In [34]: def RandomForestModel(features, target, split):
    xtrain = features[:split][linkedFeatures].reset_index(drop=True)
    ytrain = target[:split][['use [kW]']].reset_index(drop=True)
    ytest = target[split:][['use [kW]']].reset_index(drop=True)
    xtest = features[split : (split + len(ytest))][linkedFeatures].reset_in

    randomforest = RandomForestRegressor(random_state=3)
    params = dict(max_depth=[1,2,3,4,5,6,7,8,9,10])
    scoring_fnc = make_scorer(performance_metric)
    grid = GridSearchCV(randomforest,params,scoring=scoring_fnc)
    grid = grid.fit(xtrain, ytrain.values.ravel())
    model = grid.best_estimator_
    mae = mean_absolute_error(ytest, model.predict(xtest))
    return mae
```

```
In [35]: models.append("Random Forest")

maeVal = RandomForestModel(weather_hourly_B, energy_hourly_B, splitHB)
maeBHourly.append(maeVal)
print("The Mean Absolute Error for Random Forest for hourly points on House

maeVal = RandomForestModel(weather_hourly_C, energy_hourly_C, splitHC)
maeCHourly.append(maeVal)
print("The Mean Absolute Error for Random Forest for hourly points on House

maeVal = RandomForestModel(weather_hourly_F, energy_hourly_F, splitHF)
maeFHourly.append(maeVal)
print("The Mean Absolute Error for Random Forest for hourly points on House

maeVal = RandomForestModel(weather_daily_B, energy_daily_B, splitDB)
maeBDaily.append(maeVal)
print("The Mean Absolute Error for Random Forest for daily points on House

maeVal = RandomForestModel(weather_daily_C, energy_daily_C, splitDC)
maeCDaily.append(maeVal)
print("The Mean Absolute Error for Random Forest for daily points on House

maeVal = RandomForestModel(weather_daily_F, energy_daily_F, splitDF)
maeFDaily.append(maeVal)
print("The Mean Absolute Error for Random Forest for daily points on House

The Mean Absolute Error for Random Forest for hourly points on House B is
0.480918561313664
The Mean Absolute Error for Random Forest for hourly points on House C is
21.15884825500904
The Mean Absolute Error for Random Forest for hourly points on House F is
50.39096149110985
The Mean Absolute Error for Random Forest for daily points on House B is
7.6474998072644835
The Mean Absolute Error for Random Forest for daily points on House C is
209.23926077321673
The Mean Absolute Error for Random Forest for daily points on House F is
424.12097549005034
```

ARIMA

ARIMA, short for 'Auto Regressive Integrated Moving Average' is a class of models that 'explains' a given time series based on its own past values, that is, its own lags and the lagged forecast errors, so that equation can be used to forecast future values.

The parameters of the ARIMA model are defined as follows:

- p: The number of lag observations included in the model, also called the lag order.
- d: The number of times that the raw observations are differenced, also called the degree of differencing.
- q: The size of the moving average window, also called the order of moving average.

We will be iterating over the permissible values of p and d to get the best prediction in terms of minimum MAE.

```
In [36]: def warn(*args, **kwargs):  
    pass  
    import warnings  
    warnings.warn = warn
```

```
In [37]: def testForParameters(train, test, split):  
    ypredicted=[]  
    optimalp =0  
    optimald = 0  
    leastererror = float("inf")  
    for p in range(0,10):  
        for d in range(0,2):  
            newmodel = ARIMA(train, order=(p,d,0))  
            model_fit = newmodel.fit()  
            ypredicted = model_fit.predict(start=split, end=(split+len(test)  
            error = mean_absolute_error(test, ypredicted)  
            if(error<leastererror):  
                leastererror = error  
                optimalp = p  
                optimald = d  
    return optimalp,optimald  
  
def ArimaModel(ytrain, ytest, split, p, d):  
    model = ARIMA(ytrain, order=(p,d,0))  
    model = model.fit()  
    ypredict = model.predict(start=split, end=(split+len(ytest)-1))  
    print(model.summary())  
    mae = mean_absolute_error(ytest, ypredict)  
    plt.figure(figsize=(20,8))  
    plt.plot(ytest)  
    plt.plot(ypredict, color='red')  
    plt.show()  
    return mae
```

```
In [38]: models.append("ARIMA")
for i in range(6):
    target = energyData[i][['use [kW]']].values
    split = splitPoints[i]
    train = target[:split]
    test = target[split:]
    optimalp,optimald = testForParameters(train, test, split)
    print("Best P & D Values for ", tag[i]," are :", optimalp, " ", optimald)
    mae = ArimaModel(train, test, split, optimalp, optimald)
    print ("The Mean Absolute Error for ARIMA Model for ", tag[i], "is :",
    aggregate[i].append(mae)
```

Best P & D Values for House B (Hourly) are : 6 0

ARMA Model Results

```
=====
=====
Dep. Variable:          y    No. Observations:
8016
Model:                ARMA(6, 0)    Log Likelihood          -1123
9.685
Method:                css-mle    S.D. of innovations
0.983
Date:                  Thu, 18 Mar 2021    AIC                2249
5.370
Time:                  20:52:35    BIC                2255
1.284
Sample:                0    HQIC                2251
4.507
```

```
=====
=====
              coef      std err          z      P>|z|      [0.025
0.975]
-----
const          1.3330      0.034     39.235      0.000      1.266
1.400
ar.L1.y         0.6481      0.011     58.033      0.000      0.626
0.670
ar.L2.y        -0.0008      0.013     -0.060      0.952     -0.027
0.025
ar.L3.y         0.0263      0.013      1.980      0.048      0.000
0.052
ar.L4.y         0.0290      0.013      2.184      0.029      0.003
0.055
ar.L5.y        -0.0119      0.013     -0.894      0.371     -0.038
0.014
ar.L6.y        -0.0139      0.011     -1.248      0.212     -0.036
0.008
```


Roots

```
=====
=====
              Real      Imaginary      Modulus      Frequ
ency
-----
```

```

-----
AR.1          1.6300          -0.1340j          1.6355          -0.
0131
AR.2          1.6300          +0.1340j          1.6355          0.
0131
AR.3          0.2243          -1.9906j          2.0032          -0.
2321
AR.4          0.2243          +1.9906j          2.0032          0.
2321
AR.5          -2.2808          -1.2171j          2.5853          -0.
4220
AR.6          -2.2808          +1.2171j          2.5853          0.
4220
-----
-----

```



The Mean Absolute Error for ARIMA Model for House B (Hourly) is : 0.52
15728809115998

Best P & D Values for House C (Hourly) are : 5 0

ARMA Model Results

```


=====
=====
Dep. Variable:          y    No. Observations:
7656
Model:                ARMA(5, 0)    Log Likelihood          -37
190.857
Method:                css-mle    S.D. of innovations
31.148
Date:                  Thu, 18 Mar 2021    AIC                74
395.714
Time:                  20:52:44    BIC                74
444.317
Sample:                0    HQIC                74
412.386
=====
=====

```

	coef	std err	z	P> z	[0.025
0.975]					

const	51.2534	1.582	32.393	0.000	48.152
54.355					
ar.L1.y	0.7925	0.011	69.353	0.000	0.770
0.815					
ar.L2.y	-0.0236	0.015	-1.621	0.105	-0.052
0.005					
ar.L3.y	0.0195	0.015	1.338	0.181	-0.009
0.048					
ar.L4.y	-0.0043	0.015	-0.295	0.768	-0.033
0.024					
ar.L5.y	-0.0090	0.011	-0.786	0.432	-0.031
0.013					
Roots					
=====					
=====					
	Real	Imaginary	Modulus	Fre	
quency					

AR.1	1.3188	-0.0000j	1.3188	-	
0.0000					
AR.2	2.5842	-0.0000j	2.5842	-	
0.0000					
AR.3	-0.3568	-2.9629j	2.9843	-	
0.2691					
AR.4	-0.3568	+2.9629j	2.9843		
0.2691					
AR.5	-3.6688	-0.0000j	3.6688	-	
0.5000					



The Mean Absolute Error for ARIMA Model for House C (Hourly) is : 21.534172498125194

Best P & D Values for House F (Hourly) are : 1 0

ARMA Model Results

```

=====
=====
Dep. Variable:          y      No. Observations:
7656
Model:                ARMA(1, 0)    Log Likelihood      -42
787.887
Method:                css-mle      S.D. of innovations
64.704
Date:                  Thu, 18 Mar 2021    AIC              85
581.773
Time:                  20:52:52    BIC              85
602.603
Sample:                0      HQIC              85
588.919
  
```


```

=====
=====
              coef      std err          z      P>|z|      [0.025
0.975]
-----
const          97.8637      2.252      43.465      0.000      93.451
102.277
ar.L1.y         0.6717      0.008      79.338      0.000      0.655
0.688
  
```

Roots

```

=====
=====
              Real          Imaginary      Modulus      Fre
quency
-----
AR.1          1.4889          +0.0000j      1.4889
0.0000
-----
-----
  
```



The Mean Absolute Error for ARIMA Model for House F (Hourly) is : 59.17542783489664

Best P & D Values for House B (Daily) are : 8 0

ARMA Model Results

```

=====
=====
Dep. Variable:          y      No. Observations:
335
Model:                ARMA(8, 0)      Log Likelihood      -127
7.913
Method:                css-mle      S.D. of innovations      1
0.958
Date:                  Thu, 18 Mar 2021      AIC      257
5.825
Time:                  20:52:58      BIC      261
3.966
Sample:                0      HQIC      259
1.031

```

```

=====
=====
              coef      std err          z      P>|z|      [0.025
0.975]
-----
-----
const          32.3668      4.045      8.001      0.000      24.438      4
0.296
ar.L1.y         0.3820      0.055      6.937      0.000      0.274
0.490
ar.L2.y        -0.0092      0.059     -0.157      0.875     -0.124
0.106
ar.L3.y         0.1094      0.058      1.876      0.061     -0.005
0.224
ar.L4.y         0.1143      0.059      1.949      0.051     -0.001
0.229
ar.L5.y        -0.0088      0.059     -0.151      0.880     -0.124
0.106
ar.L6.y         0.1120      0.058      1.922      0.055     -0.002
0.226
ar.L7.y         0.0850      0.059      1.449      0.147     -0.030
0.200
ar.L8.y         0.0759      0.055      1.383      0.167     -0.032
0.184

```

Roots

```


=====
=====
              Real          Imaginary      Modulus      Frequ
ency
-----
-----
AR.1          1.0426          -0.0000j      1.0426      -0.
0000
AR.2          0.8617          -0.9593j      1.2895      -0.
1335
AR.3          0.8617          +0.9593j      1.2895      0.

```

```

1335
AR.4          -0.1296          -1.3691j          1.3753          -0.
2650
AR.5          -0.1296          +1.3691j          1.3753          0.
2650
AR.6          -1.4453          -0.0000j          1.4453          -0.
5000
AR.7          -1.0903          -1.2611j          1.6670          -0.
3635
AR.8          -1.0903          +1.2611j          1.6670          0.
3635

```



The Mean Absolute Error for ARIMA Model for House B (Daily) is : 6.291542500487939

Best P & D Values for House C (Daily) are : 3 0

ARMA Model Results

```

=====
=====
Dep. Variable:          y    No. Observations:
320
Model:                ARMA(3, 0)    Log Likelihood          -2
464.075
Method:                css-mle    S.D. of innovations
534.005
Date:                  Thu, 18 Mar 2021    AIC                4
938.150
Time:                  20:53:05    BIC                4
956.991
Sample:                0    HQIC                4
945.674

```

```


=====
=====
              coef      std err          z      P>|z|      [0.025
0.975]
-----
const      1234.7748      98.865      12.490      0.000      1041.004      1
428.546

```

ar.L1.y	0.5462	0.056	9.799	0.000	0.437
0.655					
ar.L2.y	0.1188	0.063	1.880	0.060	-0.005
0.243					
ar.L3.y	0.0357	0.056	0.642	0.521	-0.073
0.145					

Roots				
=====				
=====				
	Real	Imaginary	Modulus	Frequency

AR.1	1.3102	-0.0000j	1.3102	-
0.0000				
AR.2	-2.3172	-3.9980j	4.6209	-
0.3336				
AR.3	-2.3172	+3.9980j	4.6209	-
0.3336				



The Mean Absolute Error for ARIMA Model for House C (Daily) is : 247.8907938571151

Best P & D Values for House F (Daily) are : 9 0

ARMA Model Results			
=====			
=====			
Dep. Variable:	y	No. Observations:	
320			
Model:	ARMA(9, 0)	Log Likelihood	-242
5.504			
Method:	css-mle	S.D. of innovations	47
2.369			
Date:	Thu, 18 Mar 2021	AIC	487
3.008			
Time:	20:53:13	BIC	491

4.460

Sample:

0 HQIC

488

9.561

```
=====
=====
```

	coef	std err	z	P> z	[0.025	
0.975]						


const	2317.8725	147.265	15.739	0.000	2029.239	260
6.506						
ar.L1.y	0.4337	0.056	7.771	0.000	0.324	
0.543						
ar.L2.y	-0.0247	0.061	-0.408	0.683	-0.143	
0.094						
ar.L3.y	0.0283	0.056	0.505	0.613	-0.081	
0.138						
ar.L4.y	0.0571	0.056	1.015	0.310	-0.053	
0.167						
ar.L5.y	-0.0348	0.056	-0.618	0.537	-0.145	
0.076						
ar.L6.y	0.0415	0.057	0.734	0.463	-0.069	
0.152						
ar.L7.y	0.4184	0.056	7.415	0.000	0.308	
0.529						
ar.L8.y	-0.1181	0.061	-1.923	0.054	-0.238	
0.002						
ar.L9.y	0.0278	0.057	0.491	0.623	-0.083	
0.139						

Roots

```
=====
=====
```

	Real	Imaginary	Modulus	Frequ
ency				

AR.1	-1.0545	-0.4872j	1.1616	-0.
4311				
AR.2	-1.0545	+0.4872j	1.1616	0.
4311				
AR.3	-0.2404	-1.1059j	1.1317	-0.
2841				
AR.4	-0.2404	+1.1059j	1.1317	0.
2841				
AR.5	1.0503	-0.0000j	1.0503	-0.
0000				
AR.6	0.7325	-0.8654j	1.1338	-0.
1382				
AR.7	0.7325	+0.8654j	1.1338	0.
1382				
AR.8	2.1620	-3.2782j	3.9270	-0.
1572				
AR.9	2.1620	+3.2782j	3.9270	0.
1572				



The Mean Absolute Error for ARIMA Model for House F (Daily) is : 437.3616662728573

Final Results

Below are the aggregated results of MAE achieved by different models we tested above.

```
In [39]: results = pd.DataFrame({
    'Models': models,
    'House B (Hourly)': maeBHourly,
    'House C (Hourly)': maeCHourly,
    'House F (Hourly)': maeFHourly,
    'House B (Daily)': maeBDaily,
    'House C (Daily)': maeCDaily,
    'House F (Daily)': maeFDaily,
})
print("The final results on MAEs:")
display(results)
```

The final results on MAEs:

	Models	House B (Hourly)	House C (Hourly)	House F (Hourly)	House B (Daily)	House C (Daily)	House F (Daily)
0	Naive Method	0.479529	25.249532	47.941220	6.231545	252.469670	440.971126
1	Linear Regression	0.463795	21.191972	45.505913	6.157595	266.672331	431.781231
2	Decision Tree (Without Grid Search)	0.695242	31.519694	63.341204	9.442800	333.651529	430.680348
3	Decision Tree (With Grid Search)	0.480848	21.500877	50.250791	6.228120	201.460495	424.308887
4	Random Forest	0.480919	21.158848	50.390961	7.647500	209.239261	424.120975
5	ARIMA	0.521573	21.534172	59.175428	6.291543	247.890794	437.361666

Hence we find different models performing differently for different house data sets. In particular we find the goodness of HyperParameter tuning in Decision Tree Regressor (comparing both rows 2 & 3). We have models which perform better than Naive Method viz Linear Regression performing better for all instances. Decision Tree with Grid Search performing quite better for House C Daily. Thus using different techniques and machine learning algorithms, we are able to predict energy consumption for different users with performances better than our baseline Naive model.